



# Performance Monitoring (focused on Java SE)



# Topics

- What information to monitor
- What tools to use
- What level of the software stack to monitor
  - > Operating system level
  - > JVM level
  - > Application level
- What to monitor is covered per component at a given level
  - > Example. At OS level, monitor CPU usage

# Tools For Monitoring

- Definition: An act of non-intrusively collecting or observing performance data from an operating or running application.
- What to monitor and where?
  - > Operating system: cpu utilization (kernel & user), network i/o, disk i/o, memory, processes and kernel (locks).
  - > JVM: garbage collection frequency and duration, heap usage, threads, lock contention, cpu usage
  - > Application: throughput, responsiveness

# CPU Monitoring Tools at the OS-Level

# Tools For Monitoring : OS Level

- cpu
  - > vmstat (Solaris & Linux)
  - > mpstat (Solaris)
  - > prstat (Solaris)
  - > top (Linux, prefer prstat on Solaris)
  - > Task Manager (Windows)
  - > Performance Monitor (Windows)
  - > xosview (Linux)
  - > cpubar (Solaris – Performance Tools CD)
  - > iobar (Solaris – Performance Tools CD)
  - > dtrace (Solaris)

# Tools For Monitoring : vmstat

```

huntch@ditka: ~
File Edit View Terminal Tabs Help

kthr      memory          page        disk        faults        cpu
r  b  w   swap free  re  mf pi po fr de sr m0 m1 m2 m3   in   sy   cs  us sy id
0  0  0 11532280 3322672 4 18 6 0 0 0 0 1 1 1 0 1599 491 147 1 1 98
0  0  0 11406040 3186248 0 7 0 0 0 0 0 0 0 0 0 1567 348 138 0 1 99
0  0  0 11406832 3187048 0 1 0 0 0 0 0 19 19 19 0 1737 332 153 0 1 99
0  0  0 11408128 3188344 0 1 0 2 2 0 0 0 0 0 0 1573 336 134 0 1 99
0  0  0 11407912 3188128 0 1 0 0 0 0 0 0 0 0 0 1567 331 139 0 1 99
0  0  0 11407464 3187696 2 21 0 0 0 0 0 17 17 17 0 1717 331 157 0 1 99
0  0  0 11407472 3187704 0 1 0 0 0 0 0 0 0 0 0 1563 310 149 0 1 99
0  0  0 11407520 3187744 0 2 0 0 0 0 0 0 0 0 0 1565 317 144 0 1 99
0  0  0 11407504 3187720 0 1 0 0 0 0 0 0 0 0 0 1561 308 138 0 1 99
0  0  0 11407016 3187216 32 196 0 0 0 0 0 0 0 0 0 2764 4997 1381 1 2 97
0  0  0 11396280 3176232 43 351 0 0 0 0 0 0 0 0 0 5036 14128 3772 9 4 87
0  0  0 11399760 3179872 25 65 0 0 0 0 0 8 8 8 0 1740 2450 314 16 2 82
0  0  0 11407392 3187600 0 1 0 0 0 0 0 0 0 0 0 1579 332 150 17 1 82
0  0  0 11407352 3187568 0 3 0 0 0 0 0 0 0 0 0 1572 313 138 17 1 82
0  0  0 11407312 3187536 0 1 0 0 0 0 0 0 0 0 0 1573 324 143 12 1 86
0  0  0 11407280 3187504 0 1 0 2 2 0 0 0 0 0 0 1564 316 142 0 1 99
0  0  0 11407248 3187464 0 1 0 0 0 0 0 0 0 0 0 1562 319 147 0 1 99
0  0  0 11407216 3187440 0 1 0 0 0 0 0 3 3 3 0 1593 306 140 0 1 99

kthr      memory          page        disk        faults        cpu
r  b  w   swap free  re  mf pi po fr de sr m0 m1 m2 m3   in   sy   cs  us sy id
0  0  0 11407408 3187712 0 4 0 0 0 0 0 0 0 0 0 1583 557 179 0 1 98

```

# vmstat

- Virtual Memory Statistics
- reports virtual memory statistics of process, virtual memory, disk, trap, and CPU activity.
- cpu (breakdown of percentage usage of CPU time. On multiprocessors this is an average across all processors.)
  - > us - user time
  - > sy - system time
  - > id - idle time







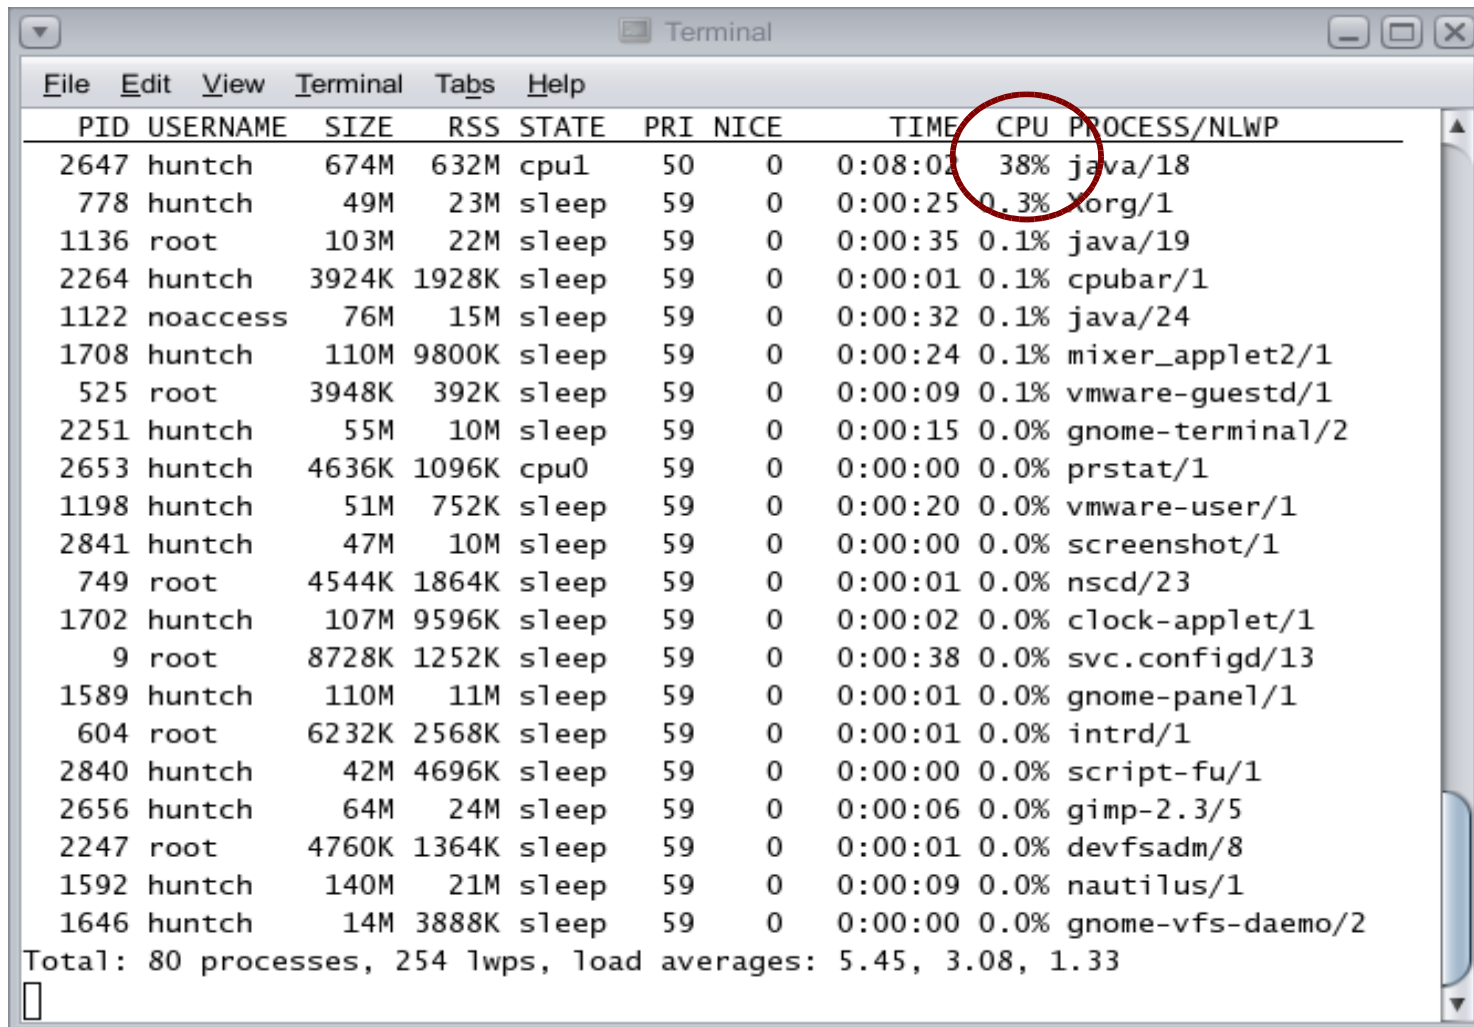
# mpstat

- Multi-processor status
- Reveals the individual CPU utilization on multi-processor
- Each row of the table represents the activity of one processor.
- The first table summarizes all activity since boot
- Each subsequent table summarizes activity for the preceding interval.
- All values are rates (events per second) unless otherwise noted.

# mpstat

- icsw involuntary context switches
- migr thread migrations (to another processor)
- usr percent user time
- sys percent system time

# Tools For Monitoring : prstat

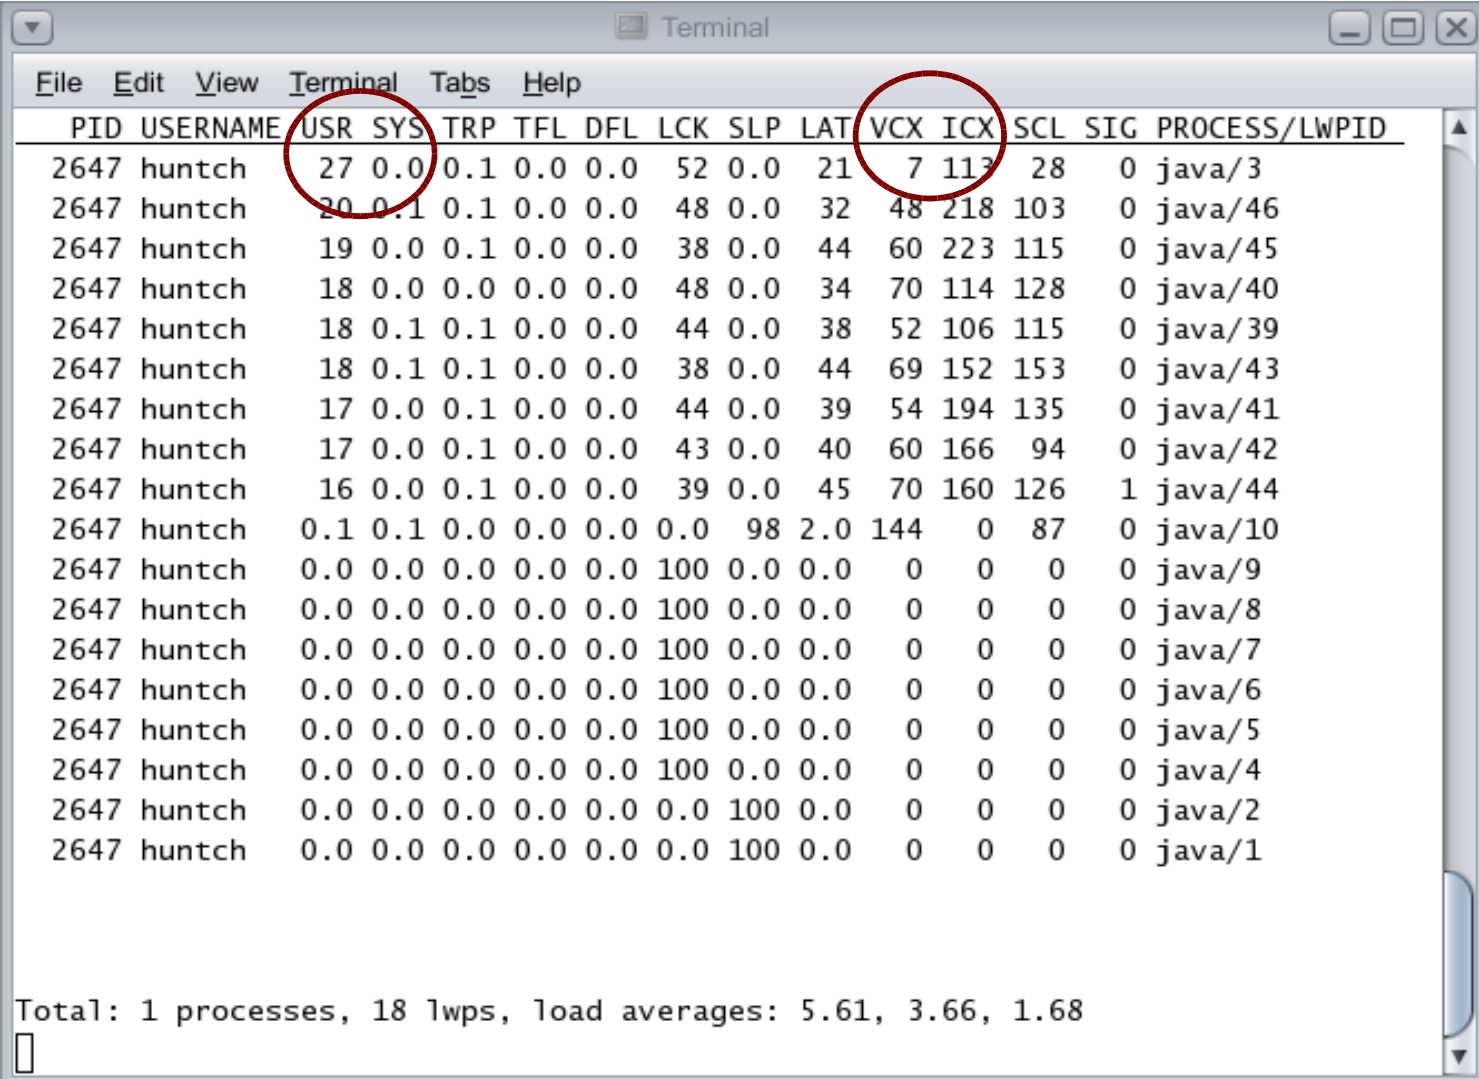


A terminal window titled "Terminal" displaying the output of the `prstat` command. The output is a table with columns: PID, USERNAME, SIZE, RSS, STATE, PRI, NICE, TIME, CPU, and PROCESS/NLWP. The first two rows are circled in red: the first row shows PID 2647 (huntch) with 674M size and 632M RSS, and the second row shows PID 778 (huntch) with 49M size and 23M RSS. The CPU column shows 38% for the first row and 0.3% for the second row. The table lists various system processes and user applications, including java, Xorg, cpubar, and gnome-terminal. At the bottom, a summary line states: "Total: 80 processes, 254 lwps, load averages: 5.45, 3.08, 1.33".

PID	USERNAME	SIZE	RSS	STATE	PRI	NICE	TIME	CPU	PROCESS/NLWP
2647	huntch	674M	632M	cpu1	50	0	0:08:01	38%	java/18
778	huntch	49M	23M	sleep	59	0	0:00:25	0.3%	Xorg/1
1136	root	103M	22M	sleep	59	0	0:00:35	0.1%	java/19
2264	huntch	3924K	1928K	sleep	59	0	0:00:01	0.1%	cpubar/1
1122	noaccess	76M	15M	sleep	59	0	0:00:32	0.1%	java/24
1708	huntch	110M	9800K	sleep	59	0	0:00:24	0.1%	mixer_applet2/1
525	root	3948K	392K	sleep	59	0	0:00:09	0.1%	vmware-guestd/1
2251	huntch	55M	10M	sleep	59	0	0:00:15	0.0%	gnome-terminal/2
2653	huntch	4636K	1096K	cpu0	59	0	0:00:00	0.0%	prstat/1
1198	huntch	51M	752K	sleep	59	0	0:00:20	0.0%	vmware-user/1
2841	huntch	47M	10M	sleep	59	0	0:00:00	0.0%	screenshot/1
749	root	4544K	1864K	sleep	59	0	0:00:01	0.0%	nscd/23
1702	huntch	107M	9596K	sleep	59	0	0:00:02	0.0%	clock-applet/1
9	root	8728K	1252K	sleep	59	0	0:00:38	0.0%	svc.configd/13
1589	huntch	110M	11M	sleep	59	0	0:00:01	0.0%	gnome-panel/1
604	root	6232K	2568K	sleep	59	0	0:00:01	0.0%	intrd/1
2840	huntch	42M	4696K	sleep	59	0	0:00:00	0.0%	script-fu/1
2656	huntch	64M	24M	sleep	59	0	0:00:06	0.0%	gimp-2.3/5
2247	root	4760K	1364K	sleep	59	0	0:00:01	0.0%	devfsadm/8
1592	huntch	140M	21M	sleep	59	0	0:00:09	0.0%	nautilus/1
1646	huntch	14M	3888K	sleep	59	0	0:00:00	0.0%	gnome-vfs-daemo/2

Total: 80 processes, 254 lwps, load averages: 5.45, 3.08, 1.33

# Tools For Monitoring : prstat -Lm



A terminal window titled "Terminal" showing the output of the `prstat -Lm` command. The output is a table with columns: PID, USERNAME, USR, SYS, TRP, TFL, DFL, LCK, SLP, LAT, VCX, ICX, SCL, SIG, and PROCESS/LWPID. The first two rows of data have red circles around the USR and SYS columns. The table lists 18 lwps for the java process. At the bottom, a summary line shows: "Total: 1 processes, 18 lwps, load averages: 5.61, 3.66, 1.68".

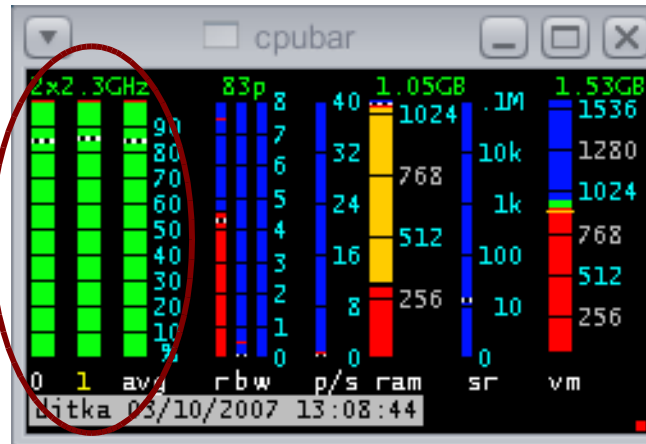
PID	USERNAME	USR	SYS	TRP	TFL	DFL	LCK	SLP	LAT	VCX	ICX	SCL	SIG	PROCESS/LWPID
2647	huntch	27	0.0	0.1	0.0	0.0	52	0.0	21	7	113	28	0	java/3
2647	huntch	20	0.1	0.1	0.0	0.0	48	0.0	32	48	218	103	0	java/46
2647	huntch	19	0.0	0.1	0.0	0.0	38	0.0	44	60	223	115	0	java/45
2647	huntch	18	0.0	0.0	0.0	0.0	48	0.0	34	70	114	128	0	java/40
2647	huntch	18	0.1	0.1	0.0	0.0	44	0.0	38	52	106	115	0	java/39
2647	huntch	18	0.1	0.1	0.0	0.0	38	0.0	44	69	152	153	0	java/43
2647	huntch	17	0.0	0.1	0.0	0.0	44	0.0	39	54	194	135	0	java/41
2647	huntch	17	0.0	0.1	0.0	0.0	43	0.0	40	60	166	94	0	java/42
2647	huntch	16	0.0	0.1	0.0	0.0	39	0.0	45	70	160	126	1	java/44
2647	huntch	0.1	0.1	0.0	0.0	0.0	0.0	98	2.0	144	0	87	0	java/10
2647	huntch	0.0	0.0	0.0	0.0	0.0	100	0.0	0.0	0	0	0	0	java/9
2647	huntch	0.0	0.0	0.0	0.0	0.0	100	0.0	0.0	0	0	0	0	java/8
2647	huntch	0.0	0.0	0.0	0.0	0.0	100	0.0	0.0	0	0	0	0	java/7
2647	huntch	0.0	0.0	0.0	0.0	0.0	100	0.0	0.0	0	0	0	0	java/6
2647	huntch	0.0	0.0	0.0	0.0	0.0	100	0.0	0.0	0	0	0	0	java/5
2647	huntch	0.0	0.0	0.0	0.0	0.0	100	0.0	0.0	0	0	0	0	java/4
2647	huntch	0.0	0.0	0.0	0.0	0.0	0.0	100	0.0	0	0	0	0	java/2
2647	huntch	0.0	0.0	0.0	0.0	0.0	0.0	100	0.0	0	0	0	0	java/1

Total: 1 processes, 18 lwps, load averages: 5.61, 3.66, 1.68

# vmstat

- CPU The percentage of recent CPU time used by the process.
- VCX The number of voluntary context switches.
- ICX The number of involuntary context switches.

# Tools For Monitoring : cpubar



- Available on Solaris Performance Tools 3.0 CD, or download from:

<http://mediacast.sun.com/share/stefanschneider/PerformanceCD3.0.tar.gz>

# CPU monitoring : What to look for

- Using HotSpot's `jps` command you can find the process ids of all running Java processes on your machine.
- Using Solaris `prstat -Lm`, or `prstat -Lmp <pid>` to locate the LWP id(s) consuming the most cpu (usr or sys) and using HotSpot's `jstack` you can find the executing threads taking the cpu (usr and sys) time.
  - > LWPID is in the far right column of `prstat`.
  - > Look for `jstack`'s corresponding 'tid', reported in hex.



# CPU monitoring : What to look for

- Idle cpu
  - > On multi-threaded applications and multi-core systems, idle cpu can be an indicator of an application's inability to scale.
  - > Combination of high sys or kernel CPU utilization and idle CPU could indicate shared resource contention as the scalability blocker.
  - > Applicable to all operating systems, i.e. Windows, Linux and Solaris

# CPU monitoring : What to look for

```

File Edit View Terminal Tabs Help
bash-3.2$ vmstat 5

```

kthr			memory		page				disk				faults				cpu				
r	b	w	swap	free	re	mf	pi	po	fr	de	sr	f0	s0	s1	s2	in	sy	cs	us	sy	id
0	0	0	659620	168400	3	19	6	2	3	0	4	-0	1	-0	1	385	886	432	3	2	95
0	0	0	129404	26456	2	87	10	0	0	0	0	0	5	0	0	2320	247894	4745	37	20	44
0	0	0	127920	25524	23	395	90	0	0	0	0	0	27	0	1	1844	241699	4280	45	21	35
0	0	0	126380	24204	1	92	0	0	0	0	0	0	0	0	0	2455	283529	4916	34	19	48
0	0	0	126380	24208	1	93	0	0	0	0	0	0	0	0	0	2469	288083	4933	33	19	49
0	0	0	126376	24204	1	79	0	0	0	0	0	0	0	0	0	2171	241732	4294	33	20	47
0	0	0	126344	24172	1	78	1	0	0	0	0	0	16	0	3	2202	242373	4782	36	20	44
0	0	0	126544	24372	1	85	0	0	0	0	0	0	0	0	0	2127	263456	4928	42	21	37
0	0	0	126504	24332	1	92	0	0	0	0	0	0	0	0	6	2498	284583	5041	33	19	48
0	0	0	126112	23940	1	144	2	0	0	0	0	0	0	0	2	2027	273269	5311	46	20	34
0	0	0	125252	22952	3	97	53	0	0	0	0	0	10	0	14	2091	231099	4423	35	20	44



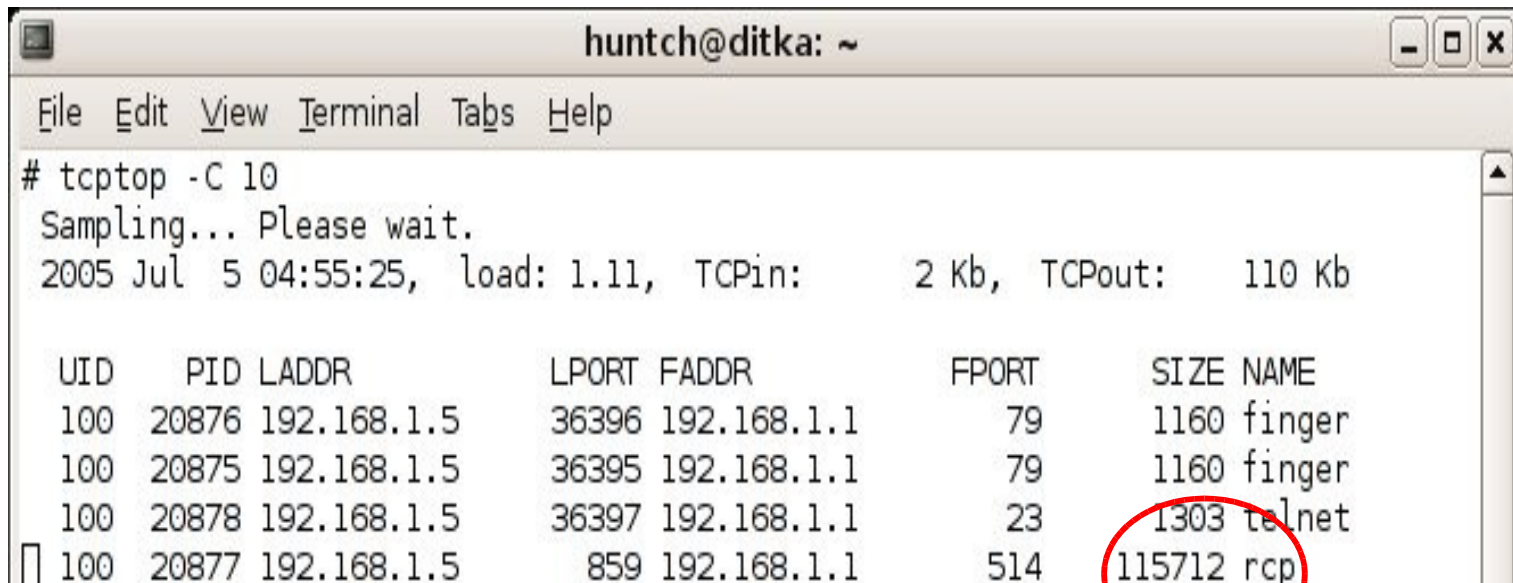
# Networking I/O Monitoring at OS Level

# Tools For Monitoring : OS Level

- network i/o
  - > netstat (Solaris & Linux)
  - > Performance Monitor (Windows)
  - > dtrace (Solaris)
  - > nicstat (Solaris – Performance Tools CD)
  - > tcptop (Dtrace Toolkit)
- Data of interest
  - > network utilization, established connections,

# Network i/o : What to look for

- tcptop can show per process tcp stats



```

# tcptop -C 10
Sampling... Please wait.
2005 Jul  5 04:55:25, load: 1.11, TCPin:      2 Kb, TCPout:    110 Kb

  UID   PID  LADDR      LPORT FADDR      FPORT   SIZE NAME
  ---   ---  ---      ---  ---      ---  ---  ---
  100   20876 192.168.1.5    36396 192.168.1.1    79    1160 finger
  100   20875 192.168.1.5    36395 192.168.1.1    79    1160 finger
  100   20878 192.168.1.5    36397 192.168.1.1    23    1303 telnet
  100   20877 192.168.1.5      859 192.168.1.1    514   115712 rcp
  
```

- Here we can see 'rcp' is generating about 115 kb of traffic

# Tools For Monitoring : nicstat

huntch@ditka: ~

File Edit View Terminal Tabs Help

huntch@ditka:~\$ nicstat 1

Time	Int	rKb/s	wKb/s	rPk/s	wPk/s	rAvs	wAvs	%Util	Sat
12:33:04	hme0	1.51	4.84	7.26	10.32	213.03	480.04	0.05	0.00
12:33:05	hme0	0.20	0.26	3.00	3.00	68.67	90.00	0.00	0.00
12:33:06	hme0	0.14	0.26	2.00	3.00	73.00	90.00	0.00	0.00
12:33:07	hme0	0.14	0.52	2.00	6.00	73.00	88.00	0.01	0.00
12:33:08	hme0	0.24	0.36	3.00	4.00	81.33	92.00	0.00	0.00
12:33:09	hme0	2.20	1.77	16.00	18.00	140.62	100.72	0.03	0.00
12:33:10	hme0	0.49	0.58	8.00	9.00	63.25	66.00	0.01	0.00
12:33:11	hme0	12.16	1830.38	185.06	1326.42	67.26	1413.06	15.09	0.00
12:33:12	hme0	19.03	3094.19	292.88	2229.11	66.53	1421.40	25.50	0.00
12:33:13	hme0	19.55	3151.87	301.00	2270.98	66.50	1421.20	25.98	0.00
12:33:14	hme0	11.99	1471.67	161.07	1081.45	76.25	1393.49	12.15	0.00
12:33:15	hme0	0.14	0.26	2.00	3.00	73.00	90.00	0.00	0.00
12:33:16	hme0	0.14	0.26	2.00	3.00	73.00	90.00	0.00	0.00
12:33:17	hme0	0.14	0.26	2.00	3.00	73.00	90.00	0.00	0.00

- Notice wAvs, write average size, during four intervals is about 1420 bytes, the MTU size.



# Disk I/O Monitoring at OS Level

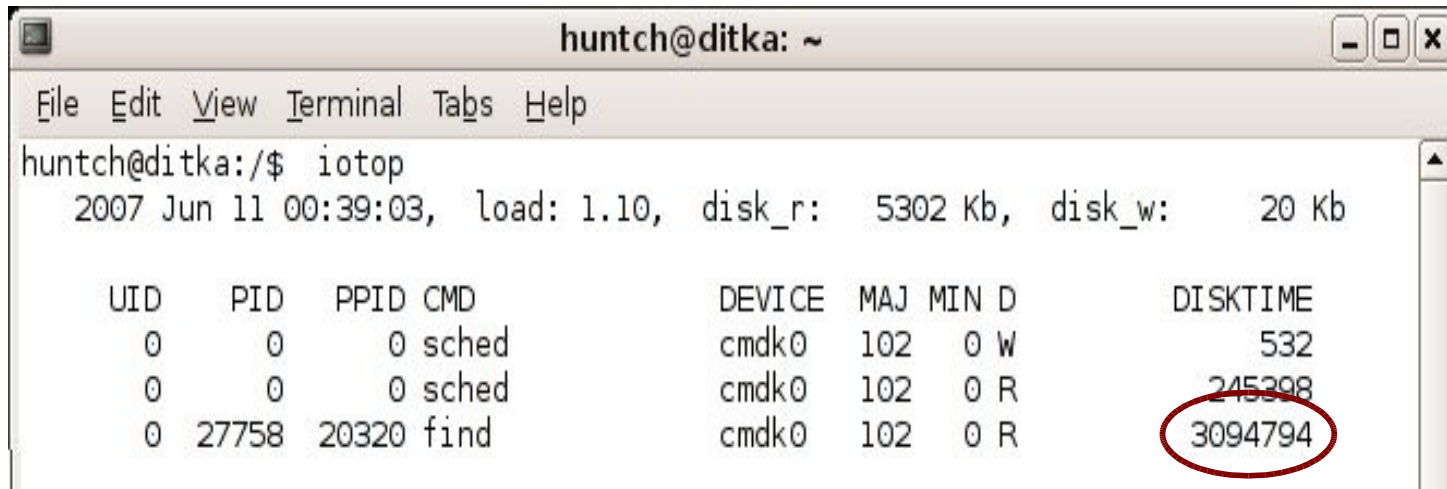
# Tools For Monitoring : OS Level

- disk i/o
  - > iostat (Solaris & Linux)
  - > Performance Monitor (Windows)
  - > dtrace (Solaris)
  - > iobar (Solaris – Performance Tools CD)
  - > iotop (Solaris – Performance Tools CD)
- Data of interest
  - > number of disk accesses, latency, average latencies

# Tools For Monitoring : iostat, iobar, iotop

- iostat reports per disk, text output
- iobar reports per disk, gui output
- iotop reports per process statistics, text output
- Data of interest
  - > number of disk accesses, latency, average latencies

# Tools For Monitoring : iotop example



```

huntch@ditka: ~
File Edit View Terminal Tabs Help
huntch@ditka:/$ iotop
2007 Jun 11 00:39:03, load: 1.10, disk_r: 5302 Kb, disk_w: 20 Kb

  UID    PID   PPID  CMD          DEVICE MAJ MIN D    DISKTIME
    0      0      0  sched        cmdk0  102  0  W      532
    0      0      0  sched        cmdk0  102  0  R      245398
    0  27758  20320  find         cmdk0  102  0  R      3094794
  
```

- iotop reporting at a 5 second interval
- DISKTIME reported in microseconds
- CMD find is keeping disk cmdk0 busy almost 60% of time during the 5 second interval

# disk i/o : What to look for

- Disk cache
  - > Why not enable disk cache?
  - > What's the risk?
  - > On Solaris x86, disk cache may be disabled by default. Linux & Windows systems usually have it enabled.
  - > Disk cache being disabled depends on the Sun branded model and how recent the model.

# Memory Monitoring at OS Level

# Tools For Monitoring : OS Level

- memory
  - > vmstat (Solaris & Linux)
  - > prstat (Solaris)
  - > top (Linux – prefer prstat on Solaris)
  - > Performance Monitor (Windows)
  - > dtrace (Solaris)
  - > cpubar (Solaris – Performance Tools CD)
  - > meminfo (Solaris – Performance Tools CD)
- Data of interest
  - > paging, memory usage



# Tools For Monitoring : vmstat

```


huntch@ditka: ~
File Edit View Terminal Tabs Help

kthr      memory          page                disk          faults          cpu
r  b  w  swap free re  mf pi po fr de sr m0 m1 m2 m3  in  sy  cs us sy id
0  0  0 11532280 3322672 4 18 6 0 0 0 0 1 1 1 0 1599 491 147 1 1 98
0  0  0 11406040 3186248 0 7 0 0 0 0 0 0 0 0 0 1567 348 138 0 1 99
0  0  0 11406832 3187048 0 1 0 0 0 0 0 19 19 19 0 1737 332 153 0 1 99
0  0  0 11408128 3188344 0 1 0 2 2 0 0 0 0 0 0 1573 336 134 0 1 99
0  0  0 11407912 3188128 0 1 0 0 0 0 0 0 0 0 0 1567 331 139 0 1 99
0  0  0 11407464 3187696 2 21 0 0 0 0 0 17 17 17 0 1717 331 157 0 1 99
0  0  0 11407472 3187704 0 1 0 0 0 0 0 0 0 0 0 1563 310 149 0 1 99
0  0  0 11407520 3187744 0 2 0 0 0 0 0 0 0 0 0 1565 317 144 0 1 99
0  0  0 11407504 3187720 0 1 0 0 0 0 0 0 0 0 0 1561 308 138 0 1 99
0  0  0 11407016 3187216 32 196 0 0 0 0 0 0 0 0 0 2764 4997 1381 1 2 97
0  0  0 11396280 3176232 43 351 0 0 0 0 0 0 0 0 0 5036 14128 3772 9 4 87
0  0  0 11399760 3179872 25 65 0 0 0 0 0 8 8 8 0 1740 2450 314 16 2 82
0  0  0 11407392 3187600 0 1 0 0 0 0 0 0 0 0 0 1579 332 150 17 1 82
0  0  0 11407352 3187568 0 3 0 0 0 0 0 0 0 0 0 1572 313 138 17 1 82
0  0  0 11407312 3187536 0 1 0 0 0 0 0 0 0 0 0 1573 324 143 12 1 86
0  0  0 11407280 3187504 0 1 0 2 2 0 0 0 0 0 0 1564 316 142 0 1 99
0  0  0 11407248 3187464 0 1 0 0 0 0 0 0 0 0 0 1562 319 147 0 1 99
0  0  0 11407216 3187440 0 1 0 0 0 0 0 3 3 3 0 1593 306 140 0 1 99

kthr      memory          page                disk          faults          cpu
r  b  w  swap free re  mf pi po fr de sr m0 m1 m2 m3  in  sy  cs us sy id
0  0  0 11407408 3187712 0 4 0 0 0 0 0 0 0 0 0 1583 557 179 0 1 98

```

# Tools For Monitoring : cpubar



- Data of interest
  - > p/s - pages per second, sr - scan rate
  - > Watch for high scan rate, or increasing trend. Low scan is ok if they are infrequent.

# memory utilization : What to look for

```

huntch@ditka: ~
File Edit View Terminal Tabs Help

huntch@ditka:~$ vmstat 5

```

kthr			memory		page				disk				faults		cpu						
r	b	w	swap	free	re	mf	pi	po	fr	de	sr	f0	s0	sl	s2	in	sy	cs	us	sy	id
1	0	0	499792	154720	1	1697	0	0	0	0	0	0	0	0	12	811	612	1761	90	7	4
1	0	0	498856	44052	1	3214	0	0	0	0	0	0	0	0	12	1290	2185	3078	66	18	15
3	0	0	501188	17212	1	1400	2	2092	4911	0	37694	0	53	0	12	5262	3387	1485	52	27	21
1	0	0	500696	20344	26	2562	18	4265	7553	0	9220	0	66	0	12	1192	3007	2733	71	17	12
1	0	0	499976	20108	3	3146	24	3032	10009	0	10971	0	63	0	6	1346	1317	3358	78	15	7
1	0	0	743664	259080	61	1706	70	8882	10017	0	19866	0	178	0	52	1213	595	688	70	12	18

- Why is swapping bad for a Java application?
- Anyone volunteers want to explain?

# memory utilization : What to look for

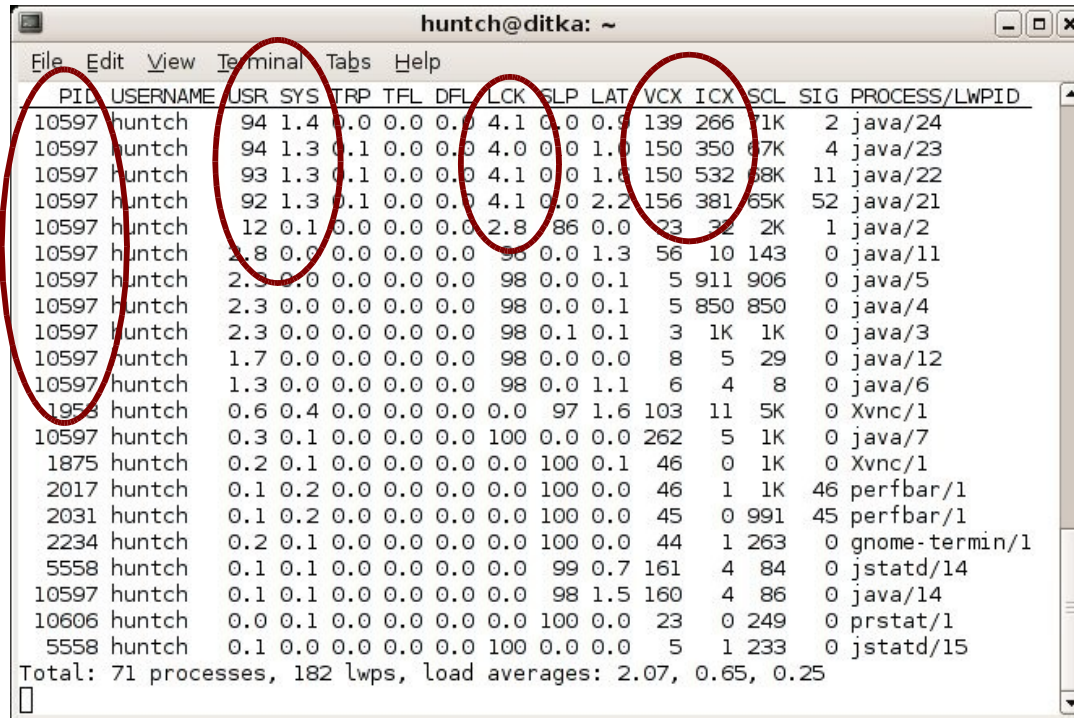
- How do you fix the swapping problem?
  - > Smaller Java heap sizes
  - > Add physical memory
  - > Reduce number of applications running on the machine
  - > Any one, or any combination of the above will help

# Processes Monitoring at OS Level

# Tools For Monitoring : OS Level

- processes
  - > ps (Solaris & Linux)
  - > vmstat (Solaris & Linux)
  - > mpstat (Solaris)
  - > prstat (Solaris)
  - > Performance Monitor (Windows)
  - > top (Linux – prefer prstat on Solaris)
  - > dtrace (Solaris)
- Data of interest
  - > footprint size, number of threads, thread state, cpu usage, runtime stack, context switches

# Tools For Monitoring : prstat -Lm



PID	USERNAME	USR	SYS	TRP	TFL	DFL	LCK	SLP	LAT	VCX	ICX	SCL	SIG	PROCESS/LWPID
10597	huntch	94	1.4	0.0	0.0	0.0	4.1	0.0	0.9	139	266	71K	2	java/24
10597	huntch	94	1.3	0.1	0.0	0.0	4.0	0.0	1.0	150	350	67K	4	java/23
10597	huntch	93	1.3	0.1	0.0	0.0	4.1	0.0	1.6	150	532	68K	11	java/22
10597	huntch	92	1.3	0.1	0.0	0.0	4.1	0.0	2.2	156	381	65K	52	java/21
10597	huntch	12	0.1	0.0	0.0	0.0	2.8	0.0	0.0	23	32	2K	1	java/2
10597	huntch	2.8	0.0	0.0	0.0	0.0	98	0.0	1.3	56	10	143	0	java/11
10597	huntch	2.3	0.0	0.0	0.0	0.0	98	0.0	0.1	5	911	906	0	java/5
10597	huntch	2.3	0.0	0.0	0.0	0.0	98	0.0	0.1	5	850	850	0	java/4
10597	huntch	2.3	0.0	0.0	0.0	0.0	98	0.1	0.1	3	1K	1K	0	java/3
10597	huntch	1.7	0.0	0.0	0.0	0.0	98	0.0	0.0	8	5	29	0	java/12
10597	huntch	1.3	0.0	0.0	0.0	0.0	98	0.0	1.1	6	4	8	0	java/6
1958	huntch	0.6	0.4	0.0	0.0	0.0	0.0	97	1.6	103	11	5K	0	Xvnc/1
10597	huntch	0.3	0.1	0.0	0.0	0.0	100	0.0	0.0	262	5	1K	0	java/7
1875	huntch	0.2	0.1	0.0	0.0	0.0	0.0	100	0.1	46	0	1K	0	Xvnc/1
2017	huntch	0.1	0.2	0.0	0.0	0.0	0.0	100	0.0	46	1	1K	46	perfbar/1
2031	huntch	0.1	0.2	0.0	0.0	0.0	0.0	100	0.0	45	0	991	45	perfbar/1
2234	huntch	0.2	0.1	0.0	0.0	0.0	0.0	100	0.0	44	1	263	0	gnome-termin/1
5558	huntch	0.1	0.1	0.0	0.0	0.0	0.0	99	0.7	161	4	84	0	jstatd/14
10597	huntch	0.1	0.1	0.0	0.0	0.0	0.0	98	1.5	160	4	86	0	java/14
10606	huntch	0.0	0.1	0.0	0.0	0.0	0.0	100	0.0	23	0	249	0	prstat/1
5558	huntch	0.1	0.0	0.0	0.0	0.0	100	0.0	0.0	5	1	233	0	jstatd/15
Total: 71 processes, 182 lwps, load averages: 2.07, 0.65, 0.25														

- Data of interest
  - > number of threads, cpu usage, locks, context switches



# Tools For Monitoring : mpstat

huntch@ditka: ~

File Edit View Terminal Tabs Help

huntch@ditka:~\$ mpstat 5

CPU	minf	mjf	xcal	intr	ithr	csw	icsw	migr	smtx	srw	syscl	usr	sys	wt	idl
0	21	1	6	319	133	261	18	23	3	0	551	4	3	0	92
1	19	1	4	36	14	59	15	23	3	0	491	4	3	0	93
2	17	1	4	19	12	61	13	23	3	0	355	4	3	0	90
3	18	1	4	16	15	49	12	23	3	0	390	4	3	0	91
CPU	minf	mjf	xcal	intr	ithr	csw	icsw	migr	smtx	srw	syscl	usr	sys	wt	idl
0	28	2	0	192	83	92	32	14	2	0	185	78	15	0	7
1	49	1	0	37	1	80	28	16	2	0	139	80	16	0	4
2	28	1	0	20	7	94	34	17	1	0	283	83	12	0	5
3	39	1	2	52	1	99	36	16	3	0	219	74	19	0	7
CPU	minf	mjf	xcal	intr	ithr	csw	icsw	migr	smtx	srw	syscl	usr	sys	wt	idl
0	34	0	2	171	75	78	32	12	1	0	173	90	9	0	2
1	38	1	0	39	1	84	29	13	2	0	153	66	12	0	23
2	28	8	0	21	9	97	31	20	2	0	167	67	13	0	20
3	35	3	1	43	1	98	29	20	3	0	190	52	25	0	23

- Data of interest
  - > context switches, lock contention

## processes : What to look for

- Why footprint size, number of threads, thread state, lock contention and context switching are important to monitor?
- What does lock contention and/or context switching look like on Solaris?
- How can you find the lock or locks causing problems?
- How can you address the thread context switching problem?

# Kernel Monitoring at OS Level

# Tools For Monitoring : OS Level

- kernel
  - > vmstat (Linux & Solaris)
  - > mpstat (Solaris)
  - > lockstat & plockstat (Solaris)
  - > Performance Monitor (Windows)
  - > dtrace (Solaris)
  - > intrstat (Solaris)
- Data of interest
  - > kernel cpu utilization, locks, system calls, interrupts, migrations, run queue depth

# Tools For Monitoring : vmstat

```
bash-3.2$ vmstat 5
```

kthr			memory		page				disk				faults		cpu						
r	b	w	swap	free	re	mf	pi	po	fr	de	sr	f0	s0	s1	s2	in	sy	cs	us	sy	id
0	0	0	659620	168400	3	19	6	2	3	0	4	-0	1	-0	1	385	886	432	3	2	95
0	0	0	129404	26456	2	87	10	0	0	0	0	0	5	0	0	2320	247894	4745	87	20	44
0	0	0	127920	25524	23	395	90	0	0	0	0	0	27	0	1	1844	241699	4280	45	21	35
0	0	0	126380	24204	1	92	0	0	0	0	0	0	0	0	0	2455	283529	4916	34	19	48
0	0	0	126380	24208	1	93	0	0	0	0	0	0	0	0	0	2469	288083	4933	33	19	49
0	0	0	126376	24204	1	79	0	0	0	0	0	0	0	0	0	2171	241732	4294	33	20	47
0	0	0	126344	24172	1	78	1	0	0	0	0	0	16	0	3	2202	242373	4782	36	20	44
0	0	0	126544	24372	1	85	0	0	0	0	0	0	0	0	0	2127	263456	4928	42	21	37
0	0	0	126504	24332	1	92	0	0	0	0	0	0	0	0	6	2498	284583	5041	33	19	48
0	0	0	126112	23940	1	144	2	0	0	0	0	0	0	0	2	2027	273269	5311	46	20	34
0	0	0	125252	22952	3	97	53	0	0	0	0	0	10	0	14	2091	231099	4423	35	20	44

- Data of interest
  - > kernel cpu utilization, run queue depth

# Tools For Monitoring : mpstat

huntch@ditka: ~

File Edit View Terminal Tabs Help

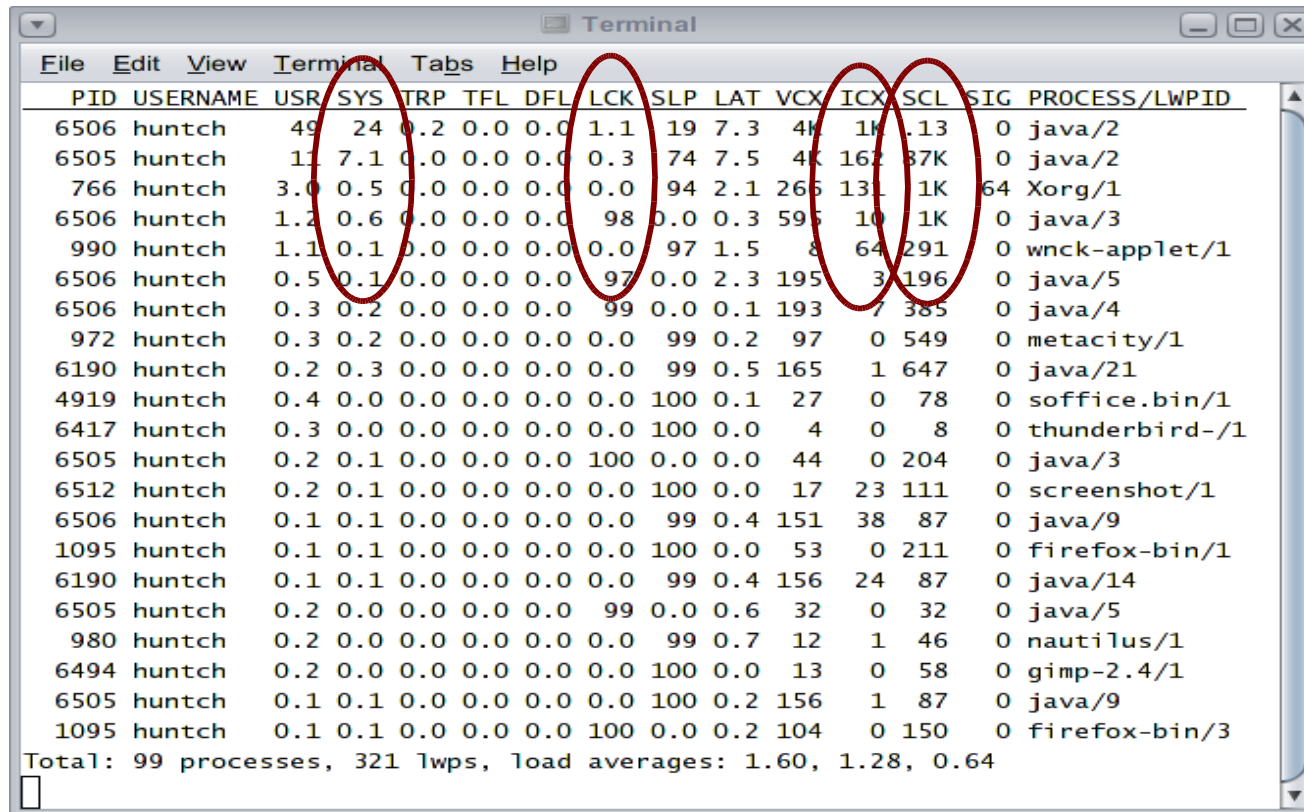
huntch@ditka:~\$ mpstat 5

CPU	minf	mjf	xcal	intr	ithr	csw	icsw	migr	smtx	srw	syscl	usr	sys	wt	idl
0	21	1	6	319	133	261	18	23	3	0	551	4	3	0	92
1	19	1	4	36	14	59	15	23	3	0	491	4	3	0	93
2	17	1	4	19	12	61	13	23	3	0	355	4	3	0	90
3	18	1	4	16	15	49	12	23	3	0	390	4	3	0	91
CPU	minf	mjf	xcal	intr	ithr	csw	icsw	migr	smtx	srw	syscl	usr	sys	wt	idl
0	28	2	0	192	83	92	32	14	2	0	185	78	15	0	7
1	49	1	0	37	1	80	28	16	2	0	139	80	16	0	4
2	28	1	0	20	7	94	34	17	1	0	283	83	12	0	5
3	39	1	2	52	1	99	36	16	3	0	219	74	19	0	7
CPU	minf	mjf	xcal	intr	ithr	csw	icsw	migr	smtx	srw	syscl	usr	sys	wt	idl
0	34	0	2	171	75	78	32	12	1	0	173	90	9	0	2
1	38	1	0	39	1	84	29	13	2	0	153	66	12	0	23
2	28	8	0	21	9	97	31	20	2	0	167	67	13	0	20
3	35	3	1	43	1	98	29	20	3	0	190	52	25	0	23

- Data of interest
  - > kernel cpu utilization, locks, system calls, interrupts, migrations



# Tools For Monitoring : prstat -Lm



PID	USERNAME	USR	SYS	TRP	TFL	DFL	LCK	SLP	LAT	VCX	ICX	SCL	SIG	PROCESS/LWPID
6506	huntch	49	24	0.2	0.0	0.0	1.1	19	7.3	4K	1K	13	0	java/2
6505	huntch	11	7.1	0.0	0.0	0.0	0.3	74	7.5	4K	162	87K	0	java/2
766	huntch	3.0	0.5	0.0	0.0	0.0	0.0	94	2.1	265	131	1K	64	Xorg/1
6506	huntch	1.2	0.6	0.0	0.0	0.0	0.0	98	0.0	0.3	595	10	0	java/3
990	huntch	1.1	0.1	0.0	0.0	0.0	0.0	97	1.5	8	64	291	0	wnck-applet/1
6506	huntch	0.5	0.1	0.0	0.0	0.0	0.0	97	0.0	2.3	195	3	0	java/5
6506	huntch	0.3	0.2	0.0	0.0	0.0	0.0	99	0.0	0.1	193	7	0	java/4
972	huntch	0.3	0.2	0.0	0.0	0.0	0.0	99	0.2	97	0	549	0	metacity/1
6190	huntch	0.2	0.3	0.0	0.0	0.0	0.0	99	0.5	165	1	647	0	java/21
4919	huntch	0.4	0.0	0.0	0.0	0.0	0.0	100	0.1	27	0	78	0	soffice.bin/1
6417	huntch	0.3	0.0	0.0	0.0	0.0	0.0	100	0.0	4	0	8	0	thunderbird-/1
6505	huntch	0.2	0.1	0.0	0.0	0.0	100	0.0	0.0	44	0	204	0	java/3
6512	huntch	0.2	0.1	0.0	0.0	0.0	0.0	100	0.0	17	23	111	0	screenshot/1
6506	huntch	0.1	0.1	0.0	0.0	0.0	0.0	99	0.4	151	38	87	0	java/9
1095	huntch	0.1	0.1	0.0	0.0	0.0	0.0	100	0.0	53	0	211	0	firefox-bin/1
6190	huntch	0.1	0.1	0.0	0.0	0.0	0.0	99	0.4	156	24	87	0	java/14
6505	huntch	0.2	0.0	0.0	0.0	0.0	99	0.0	0.6	32	0	32	0	java/5
980	huntch	0.2	0.0	0.0	0.0	0.0	0.0	99	0.7	12	1	46	0	nautilus/1
6494	huntch	0.2	0.0	0.0	0.0	0.0	0.0	100	0.0	13	0	58	0	gimp-2.4/1
6505	huntch	0.1	0.1	0.0	0.0	0.0	0.0	100	0.2	156	1	87	0	java/9
1095	huntch	0.1	0.1	0.0	0.0	0.0	100	0.0	0.2	104	0	150	0	firefox-bin/3

Total: 99 processes, 321 lwps, load averages: 1.60, 1.28, 0.64

- Data of interest
  - > kernel cpu utilization, locks, system calls, involuntary context switches

# kernel : What to look for

- Why are high sys / kernel cpu, run queue depth, lock contention, migrations and context switching are important to monitor?
  - > Discussion
- What do they indicate when each is observed?
  - > Discussion
- How to do you address each of these problems?
  - > Discussion



# Monitoring Tools at the JVM-Level

# Tools For Monitoring: JVM

- Garbage Collection
  - > -verbose:gc
  - > -XX:+PrintGCTimeStamps
  - > -XX:+PrintGCDetails
  - > -XX:+PrintGCApplicationStoppedTime
  - > -XX:+PrintGCApplicationConcurrentTime
  - > jstat, jps
  - > Jconsole
  - > VisualVM
  - > VisualGC
  - > dtrace (HotSpot JDK 6 contains samples)

# Tools For Monitoring: JVM

- Garbage Collection Data of Interest
  - > Frequency and duration of collections
  - > Java heap usage
  - > Number of application threads
  - > Lock contention
  - > CPU usage

# Tools For Monitoring : `-verbose:gc`

- `-verbose:gc`
  - > [GC 1884K->1299K(5056K), 0.0031820 secs]
- With `-XX:+PrintGCTimeStamps`
  - > 3.791: [GC 1884K->1299K(5056K), 0.0031820 secs]
- Data of interest
  - > frequency and duration, heap usage
- Explain what pattern(s) indicate potential problems.
  - > Quick demo

# Tools For Monitoring : GCDetails

- `-XX:+PrintGCDetails`
- `[GC [DefNew: 490K->64K(960K), 0.0032800 secs] 5470K->5151K(7884K), 0.0033270 secs] [Times: user=0.00 sys=0.00, real=0.00 secs]`
- `[Full GC (System) [Tenured: 5087K->5151K(6924K), 0.0971070 secs] 6047K->5151K(7884K), [Perm : 11178K->11178K(16384K)], 0.0972120 secs] [Times: user=0.10 sys=0.01, real=0.10 secs]`
- Data of interest
  - > frequency and duration, heap usage

# Tools For Monitoring : pause time

- `-XX:+PrintGCApplicationStoppedTime`
- `-XX:+PrintGCApplicationConcurrentTime`
- Helpful when tuning pause time sensitive applications
- Useful for identifying odd pause time issues when combined with gc timestamps and gc duration.

# Tools For Monitoring : jps

- jps
  - > included in the HotSpot JDK.
  - > command line utility to find running java processes.
  - > jps [-q] [-mlvV] [<hostid>] where <hostid> = <hostname>[:<port>]

# Tools For Monitoring : jstat

- jstat
  - > included in the HotSpot JDK.
  - > command line utility.
  - > `jstat -<option> [-t] [-h<lines>] <vmid> [<internal> [<count>]]`
  - > Garbage collection option(s):
    - `-gc, -gccapacity, -gccause, -gcnew, -gcnewcapacity, -gcold, -gcoldcapacity, -gcpermcapacity, -gcutil`



# Tools For Monitoring : jstat

- *Beware:* When using the concurrent collector (CMS), jstat reports two full gc events per CMS cycle, obviously misleading.
- But, young generation stats are accurate with CMS.

# Tools For Monitoring : jconsole

- jconsole
  - > Monitoring and management GUI console.
  - > Included in the HotSpot JDK.
  - > Attach local or remote.
  - > Monitor internals of a target JVM.
  - > Monitor multiple JVMs.
  - > Explain what patterns indicate potential problems.

# Tools For Monitoring : jconsole

- Endless observability
  - > MBean support for
    - JVM memory usage by memory pool / spaces
    - Class loading, JIT compilation, garbage collector, runtime, threading and logging
    - Thread monitor contention
  - > Graphical view of heap memory, threads, cpu usage and class loading

# Tools For Monitoring : VisualVM

- VisualVM
  - > Open source project at <https://visualvm.dev.java.net>
  - > Integrates several existing JDK software tools, lightweight memory and CPU profiling capabilities.
    - JConsole
    - Subset of NetBeans Profiler
  - > Includes performance analysis and troubleshooting abilities.
    - Thread deadlock detection
    - Thread monitor contention

# Tools For Monitoring : VisualVM

- VisualVM
  - > Can be further extended with specific functionality for target application via additional plug-in or extending an existing plug-in.
    - Possibilities include; GlassFish performance monitoring plug-in, JavaDB performance monitoring plug-in and external vendors such as WebSphere performance monitoring plug-in.
  - > Plugins, enhancements and updates delivered through VisualVM plug-in center.

# Tools For Monitoring : VisualVM

- VisualVM
  - > Explain what patterns indicate potential performance issues.

# Tools For Monitoring : VisualGC

- VisualGC
  - > Standalone GUI or VisualVM plug-in.
  - > Not included in HotSpot JDK. Separate download.
  - > Visually observe garbage collection behavior. (A picture is worth a thousand words).
  - > Also includes classloading and JIT compilation information.

# Tools For Monitoring : JVM

- JIT Compilation
  - > jstat
  - > Jconsole
  - > VisualVM
  - > VisualGC
  - > -XX:+PrintCompilation (can be intrusive)
  - > -XX:+LogCompilation (can be intrusive)
  - > DTrace (HotSpot JDK 6 contains samples)
- Data of interest
  - > frequency, duration, possible opt / de-opt cycles, failed compilations



# Tools For Monitoring : JIT

- **-XX:+PrintCompilation**

- 1 java.util.Properties\$LineReader::readLine (452 bytes)
- 2 java.lang.String::hashCode (60 bytes)
- 3 java.lang.String::equals (88 bytes)
- 3 made not entrant (2) java.lang.String::equals (88 bytes)
- 4 java.lang.Object::<init> (1 bytes)
- 5 java.lang.String::indexOf (151 bytes)

- **Data of interest**

- > frequency, duration, possible opt / de-opt cycles, (explain the patterns which indicate trouble)

# Tools For Monitoring : JIT

- -XX:+LogCompilation
  - > Beware, it can be intrusive
- Will probably need someone from JIT compiler team to analyze it.
- Data of interest
  - > frequency, duration, possible opt / de-opt cycles

# Tools For Monitoring : JIT

- Using .hotspot\_compiler file
- When & why to use it?
  - > JIT compiler in an endless loop attempting a “heroic” optimization which will not converge
  - > JIT compiler in a de-optimization – re-optimization cycle
  - > JIT compiler producing 'bad' code resulting in a core dump or other severe problem

# Tools For Monitoring : JIT

- Using .hotspot\_compiler file, continued ...
- What is the format?
  - > exclude A/B/C/D methodName where
  - > A.B.C.D is the fully qualified package and class name and methodName is the method name.
  - > To exclude java.util.HashMap.clear(), specify:
    - exclude java/util/HashMap clear

# Monitoring Tools at the Application-Level

# Tools For Monitoring : Application

- Application throughput and / or responsiveness
  - > JConsole (application Mbeans)
  - > Extend VisualVM with a plug-in to gather Java application data of interest and monitor the application with VisualVM
  - > Application log
  - > Specialized DTrace scripts
- Data of interest
  - > critical application information and instrumentation



# Performance Monitoring (focused on Java SE)

